

EMPLOYMENT

LexisNexis Legal & Professional

Senior Data Scientist II

Raleigh, NC

Oct 2026 - Present

- Benchmarked 15 open-source structured decision models (System One) against four LLM baselines across six classification tasks; fine-tuned three Pareto-optimal models (best-model held-out macro-F1: 92.7%) to improve latency (M4 Pro, median: 177 ms, p95: 254 ms).

Senior Data Scientist I

May 2024 - Oct 2026

- Technical lead for an agentic legal drafting system; led its migration from RAG and continue to own development.
- On a 3-person team, increased the share of drafts rated suitable for users by attorney reviewers from 40% to 86% for complaints and 59.2% to 76% for motions.
- Reduced customer-facing product runtime by 24.8% through LLM call profiling and parallelization.
- Built a contrastive-learning filter using InfoNCE with alignment and uniformity losses to remove unrelated legal passages from search results; median filtering latency of 12 ms across 40 queries.
- Built team developer tooling: a Rust MITM proxy with request replay, a web UI for visualizing intermediate agent steps, and custom Python lint rules beyond `ruff`'s rule set.
- Mentored an intern on a large-scale, full factorial experiment investigating Pareto-optimal configurations for draft quality and latency.
- Profiled high pod memory usage and introduced lazy module loading, reducing time at peak memory from 86.4% to 68% of runtime.
- Technology:** Python, Litestar, Google ADK, React, Rust, TypeScript, Tailwind

North Carolina State University

PhD Student

Raleigh, NC

Aug 2019 - Jul 2024

- Hyper-parameter optimization:** Proposed an optimizer that outperformed prior methods and ran 36.4% faster than Hyperopt on datasets used in the NeurIPS 2020 BBO Challenge.
- Defect prediction:** Proposed GHOST, an adaptive oversampling and hyper-parameter tuning pipeline that outperformed prior deep-learning methods on 14 of 20 defect-prediction datasets.
- Semi-supervised learning:** Developed GHOST2 for identifying actionable static-analysis warnings; achieved median AUC of 0.92 across eight open-source Java projects using 10% of the training labels.
- Teaching:** TA for 830 students over 5 semesters across CSC 230 (C and Software Tools), CSC 510 (Software Engineering), and CSC 591/791 (Automated Software Engineering)

RELEVANT PROJECTS

pi-permission-system (maintained fork)

TypeScript, Node.js

May 2026 - Present

[GitHub](#)

Maintain a permission-system fork for the Pi coding agent. Redesigned shell permissions using grammar-based analysis of pipelines, substitutions, and redirections, with protected-path checks and argument-aware approval rules. Extended enforcement to background commands and fixed concurrent subagent approval handling.

Q&A System for Zotero

Rust, SQLite, LanceDB

Oct 2024 - Present

[GitHub](#) :: [crates.io](#)

Built an academic question-answering and semantic search system over a user's Zotero library. Developed a Rust RAG library with configurable LLM, embedding, and reranking providers, tool calling, and limits on concurrent embedding requests. Wrote a custom PDF parser for academic papers (median throughput $\approx$ 2,400 pages/s across 11 PDFs; M4 Pro).

RAISE

Python, Keras, PyTorch

Aug 2020 - Jun 2026

[GitHub](#) :: [PyPI](#)

Sole developer of a Python package for reproducible ML experiments, providing data loading, evaluation metrics, statistical tests, and reference implementations of research algorithms. Uses Keras with a PyTorch backend; tested on Python 3.11-3.14 across Linux, Windows, and Apple Silicon macOS. Used by multiple research groups; downloaded 79k+ times on PyPI.

Flatness-guided HPO

Python, PyTorch, Keras

Nov 2023 - May 2026

[GitHub](#) :: [Paper](#)

Developed a hyper-parameter optimizer using closed-form loss-curvature estimates to select configurations with flatter losses. Implemented comparative benchmarks on MNIST, SVHN, OpenML, and Bayesmark.

GHOST

Python, Keras

Jun 2020 - Mar 2022

[GitHub](#) :: [Paper](#)

Developed the GHOST reference implementation for software defect prediction, combining adaptive oversampling with DODGE hyper-parameter optimization. Implemented configurable ablation experiments and Scott-Knott statistical comparisons using RAISE.

SKILLS

Languages: Rust, Python, TypeScript, Java, C++, Gleam

ML/AI: Keras, PyTorch, Hyper-parameter optimization, RAG, LLM evaluation, Agentic systems, Google ADK

Frameworks: React, Flask, Litestar, FastAPI

Databases: LanceDB, MySQL, MongoDB, DynamoDB

Cloud: Google Cloud (Compute Engine, Cloud Storage, Vertex AI), AWS (EC2, S3, DynamoDB)

EDUCATION

North Carolina State University

Raleigh, NC

Ph.D. Computer Science

Aug 2019 - Jul 2024

- Advisor: *Dr. Tim Menzies*

- Dissertation: [Guidelines for the Application of Neural Technologies in Software Analytics \(or: How to Do More with Less in SE\)](#)

PES University

Bangalore, India

B.E. Computer Science – Advisor: Dr. Snehanshu Saha

Aug 2015 - May 2019

HONORS

Distinguished Reviewer, ICSE 2026

Mar 2026

Google Developer Expert - ML and GCP

Mar 2025 - Present

Google Cloud Champion Innovator - Cloud AI/ML

Oct 2022 - Mar 2025

Google Cloud Research Innovator

Feb 2022

RECENT PUBLICATIONS

See full list on [Google Scholar](#).

1. Nagarsekar, A., Bhat, M. P. A., Nesarkar, A., Godhwani, V., **Yedida, R.**, Challa, A., Danda, S., & Saha, S. (2026). A Generalisation Signal Need Not Be a Model-Selection Signal. In *I Can't Believe It's Not Better: Failure Modes of AI in Biology Workshop (NeurIPS 2026)*.
2. **Yedida, R.**, & Menzies, T. (2025). Is Hyper-Parameter Optimization Different for Software Analytics? *IEEE Transactions on Software Engineering*, 51(6).
3. Baldassarre, M. T., Ernst, N., Hermann, B., Menzies, T., & **Yedida, R.** (2023). (Re)use of Research Results (is Rampant). *Communications of the ACM*, 66(2), 75-81.
4. **Yedida, R.**, Kang, H. J., Tu, K., Lo, D., & Menzies, T. (2023). How to Find Actionable Static Analysis Warnings: A Case Study with FindBugs. *IEEE Transactions on Software Engineering*, 49(4), 2856-2872.
5. **Yedida, R.**, Krishna, R., Kalia, A., Menzies, T., Xiao, J., & Vukovic, M. (2023). An Expert System for Redesigning Software for Cloud Applications. *Expert Systems with Applications*, 219, 119673.